Module outline

Application:

- Text classification with ATIS data

Model:

- Recurrence
- Long-short term memory cell
- Different recurrent networks

Concept:

- Embedding

Train-Test-Predict Workflow

Sequences (many to many)

Problem: Tagging entities in Air Traffic Controller (ATIS) data

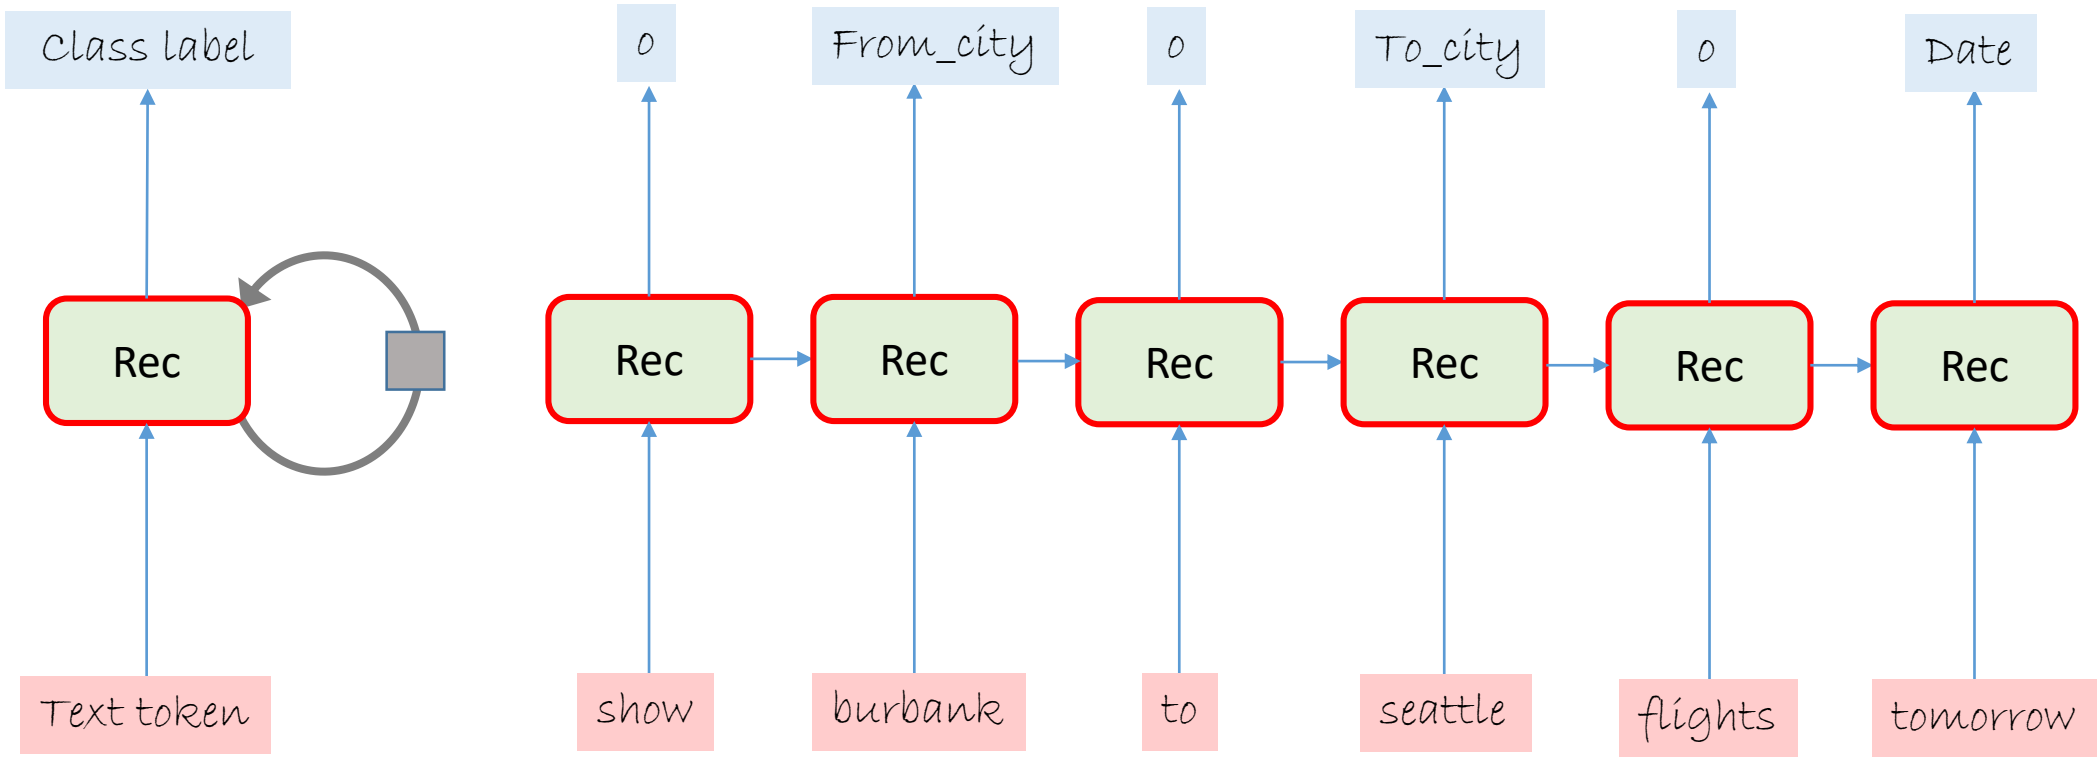

ATIS data

Domain:

✓ ATIS contains human-computer queries from the domain of Air Travel Information Services.

Data summary:

✓ 943 unique words a.k.a. : Vocabulary

✓ 129 unique tags a.k.a.: Labels

✓ 26 intent tags: not used in this tutorial

Sequence Id	Input Word (sample)	Word Index (in vocabulary) S0	Word Label	Label Index (S2)
19	# BOS	178:1	# O	128:1
19	# please	688:1	# O	128:1
19	# give	449:1	# O	128:1
19	# me	581:1	# O	128:1
19	# the	827:1	# O	128:1
19	# flights	429:1	# O	128:1
19	# from	444:1	# O	128:1
19	# boston	266:1	# B-fromloc.city_name	48:1
19	# to	851:1	# O	128:1
19	# pittsburgh	682:1	# B-toloc.city_name	78:1
19	# on	654:1	# O	128:1
19	# thursday	845:1	# B-depart_date.day_name	26:1
19	# of	646:1	# O	128:1
19	# next	621:1	# B-depart_date.date_relative	25:1
19	# week	910:1	# O	128:1
19	# EOS	179:1	# O	128:1

Sequence Id: 19 indicates – this sentence is the 19th sentence in the data set
Word Index: ###:1 indicates the position of the corresponding word in the vocabulary (total 943 words)
Label Index: ###:1 indicates the position of the corresponding tag in tag index (total 129 tags)

Sequence Tagging (Input / Label Pre-processing)

Create a numerical representation of the input words

For MNIST data:

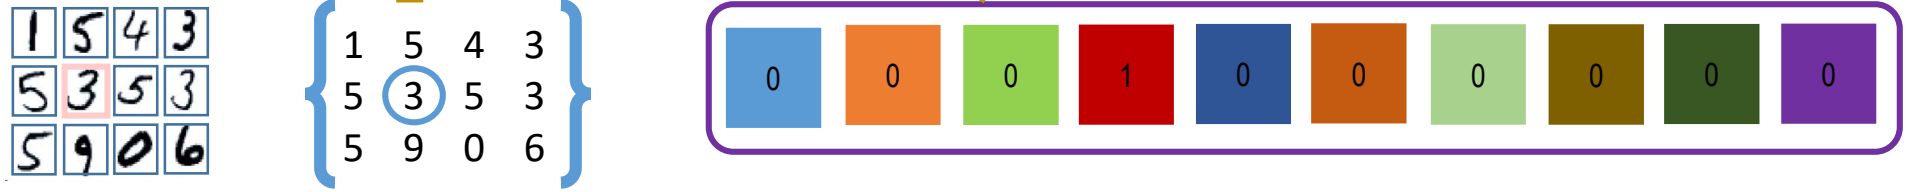

For each word - One-hot representation is a vector with 943 elements

# boston	266:1

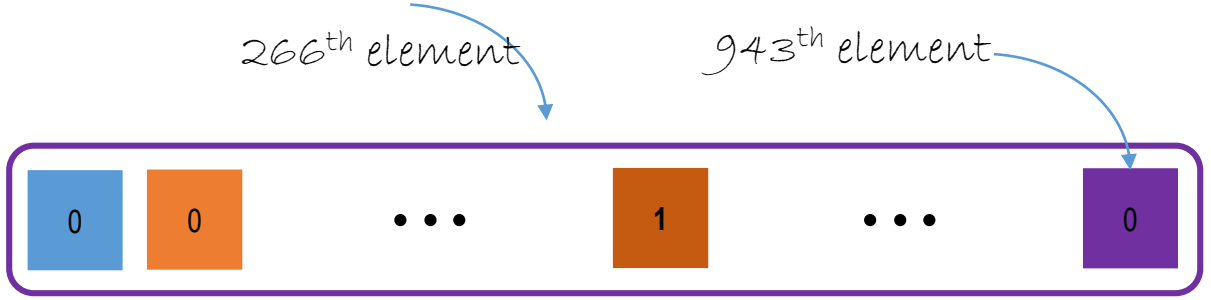

For each label – one-hot representation is a vector with 129 elements

Embedding

One-hot Encoding

Numerical representation of text

Word Embedding

Technique to map words or phrases to vector of real numbers.

Maps **one-hot encoded** vector to a lower dimensional space

Linear Embedding

Multiply a matrix with one-hot encoded vector ($\mathbf{W}_e \vec{X}^T$)

$\vec{X}^T$: vector of size 1 x 943

$\mathbf{W}_e$: matrix of size 150 x 943

Popular Embedding

GloVe (https://en.wikipedia.org/wiki/GloVe_(machine_learning))

Word2Vec (https://en.wikipedia.org/wiki/Word2vec)

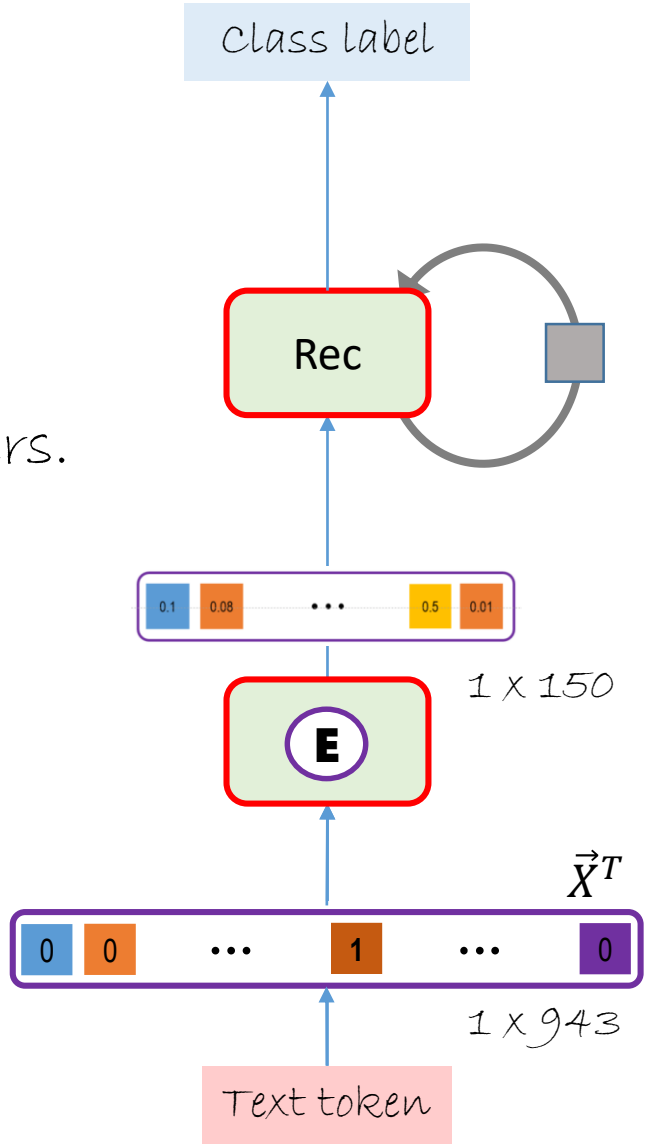

Model

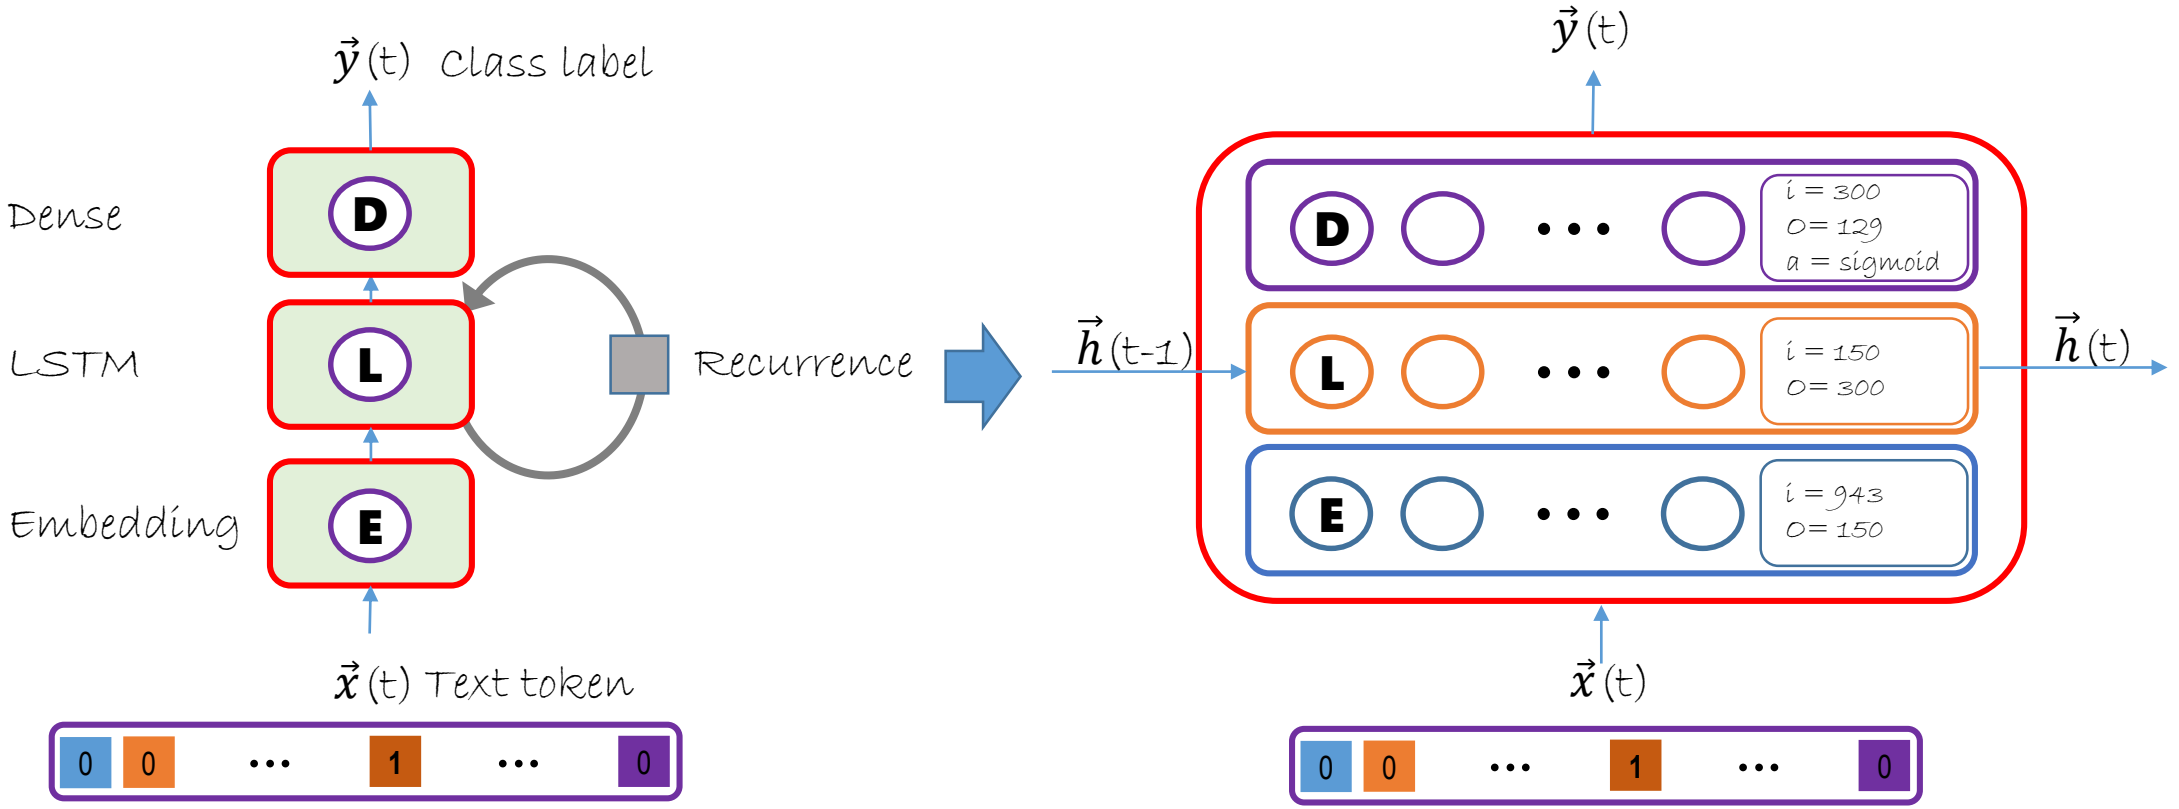
$\vec{y}$(t) Class label
Dense
D
LSTM
L
Recurrence
Embedding
E
$\vec{x}$(t) Text token
0
0
1
0
$\vec{y}$(t)
$\vec{h}$(t-1)
$\vec{h}$(t)
i = 300
o= 129
a = sigmoid
i = 150
o= 300
i = 943
o= 150
$\vec{x}$(t)

Text classification

Problem: Tagging entities in Air Traffic Controller (ATIS) data

# BOS	# O
# from	# O
# boston	# B-fromloc.city_name
# to	# O
# pittsburgh	# B-toloc.city_name
# on	# O
# thursday	# B-depart_date.day_name
# of	# O
# next	# B-depart_date.date_relative
# week	# O
# EOS	# O

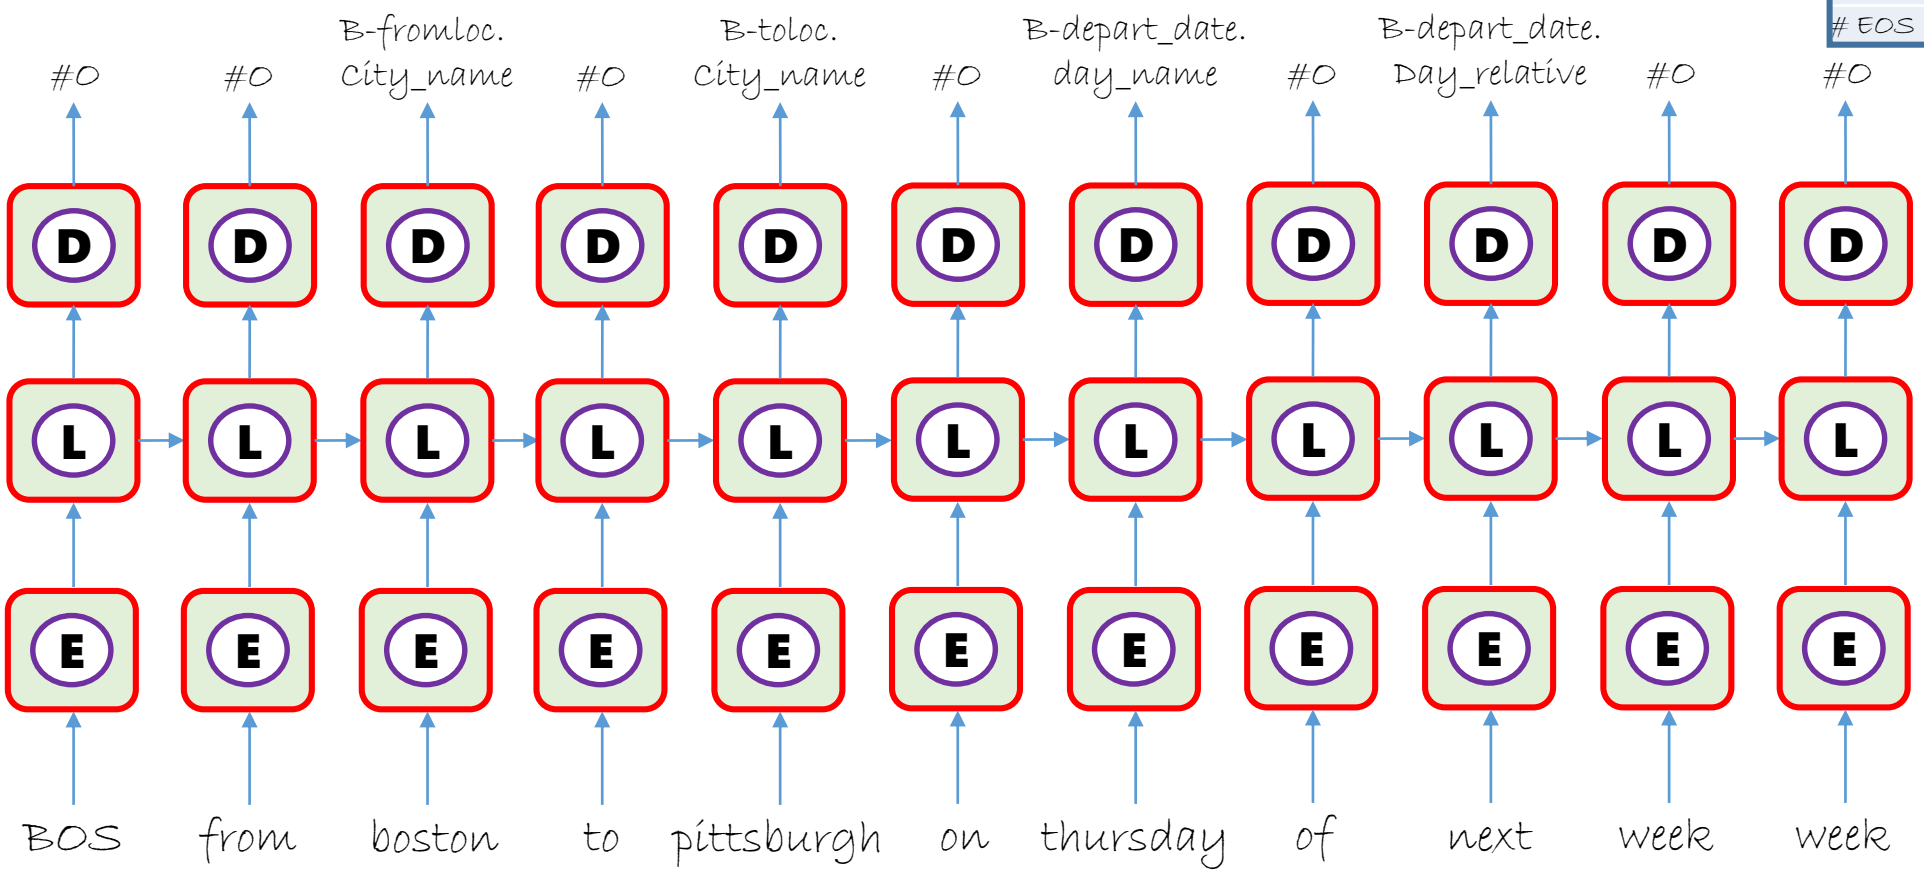

Text classification

Problem: Tagging entities in Air Traffic Controller (ATIS) data

# BOS	# O
# from	# O
# boston	# B-fromloc.city_name
# to	# O
# pittsburgh	# B-toloc.city_name
# on	# O
# thursday	# B-depart_date.day_name
# of	# O
# next	# B-depart_date.date_relative
# week	# O
# EOS	# O

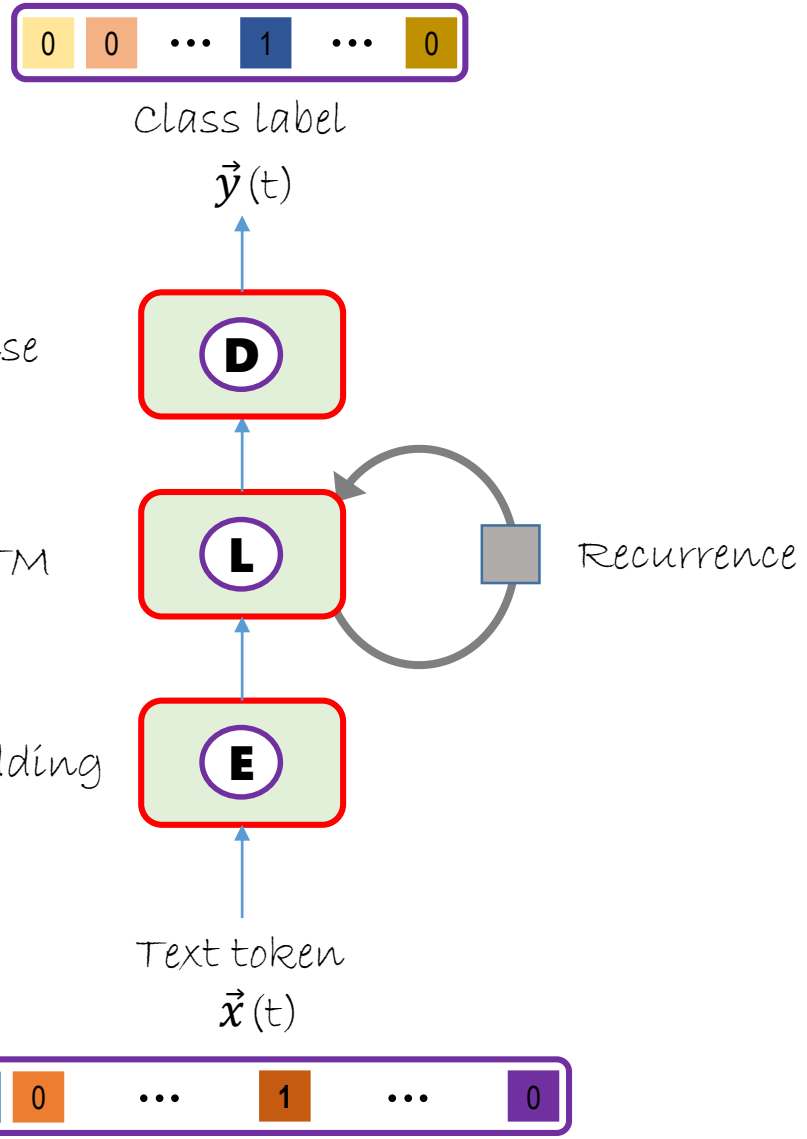

'BOS from boston to Pittsburgh on Thursday of next week EOS'

Input feature (1 x 11 x (1x943))

#1

Error or Loss Function

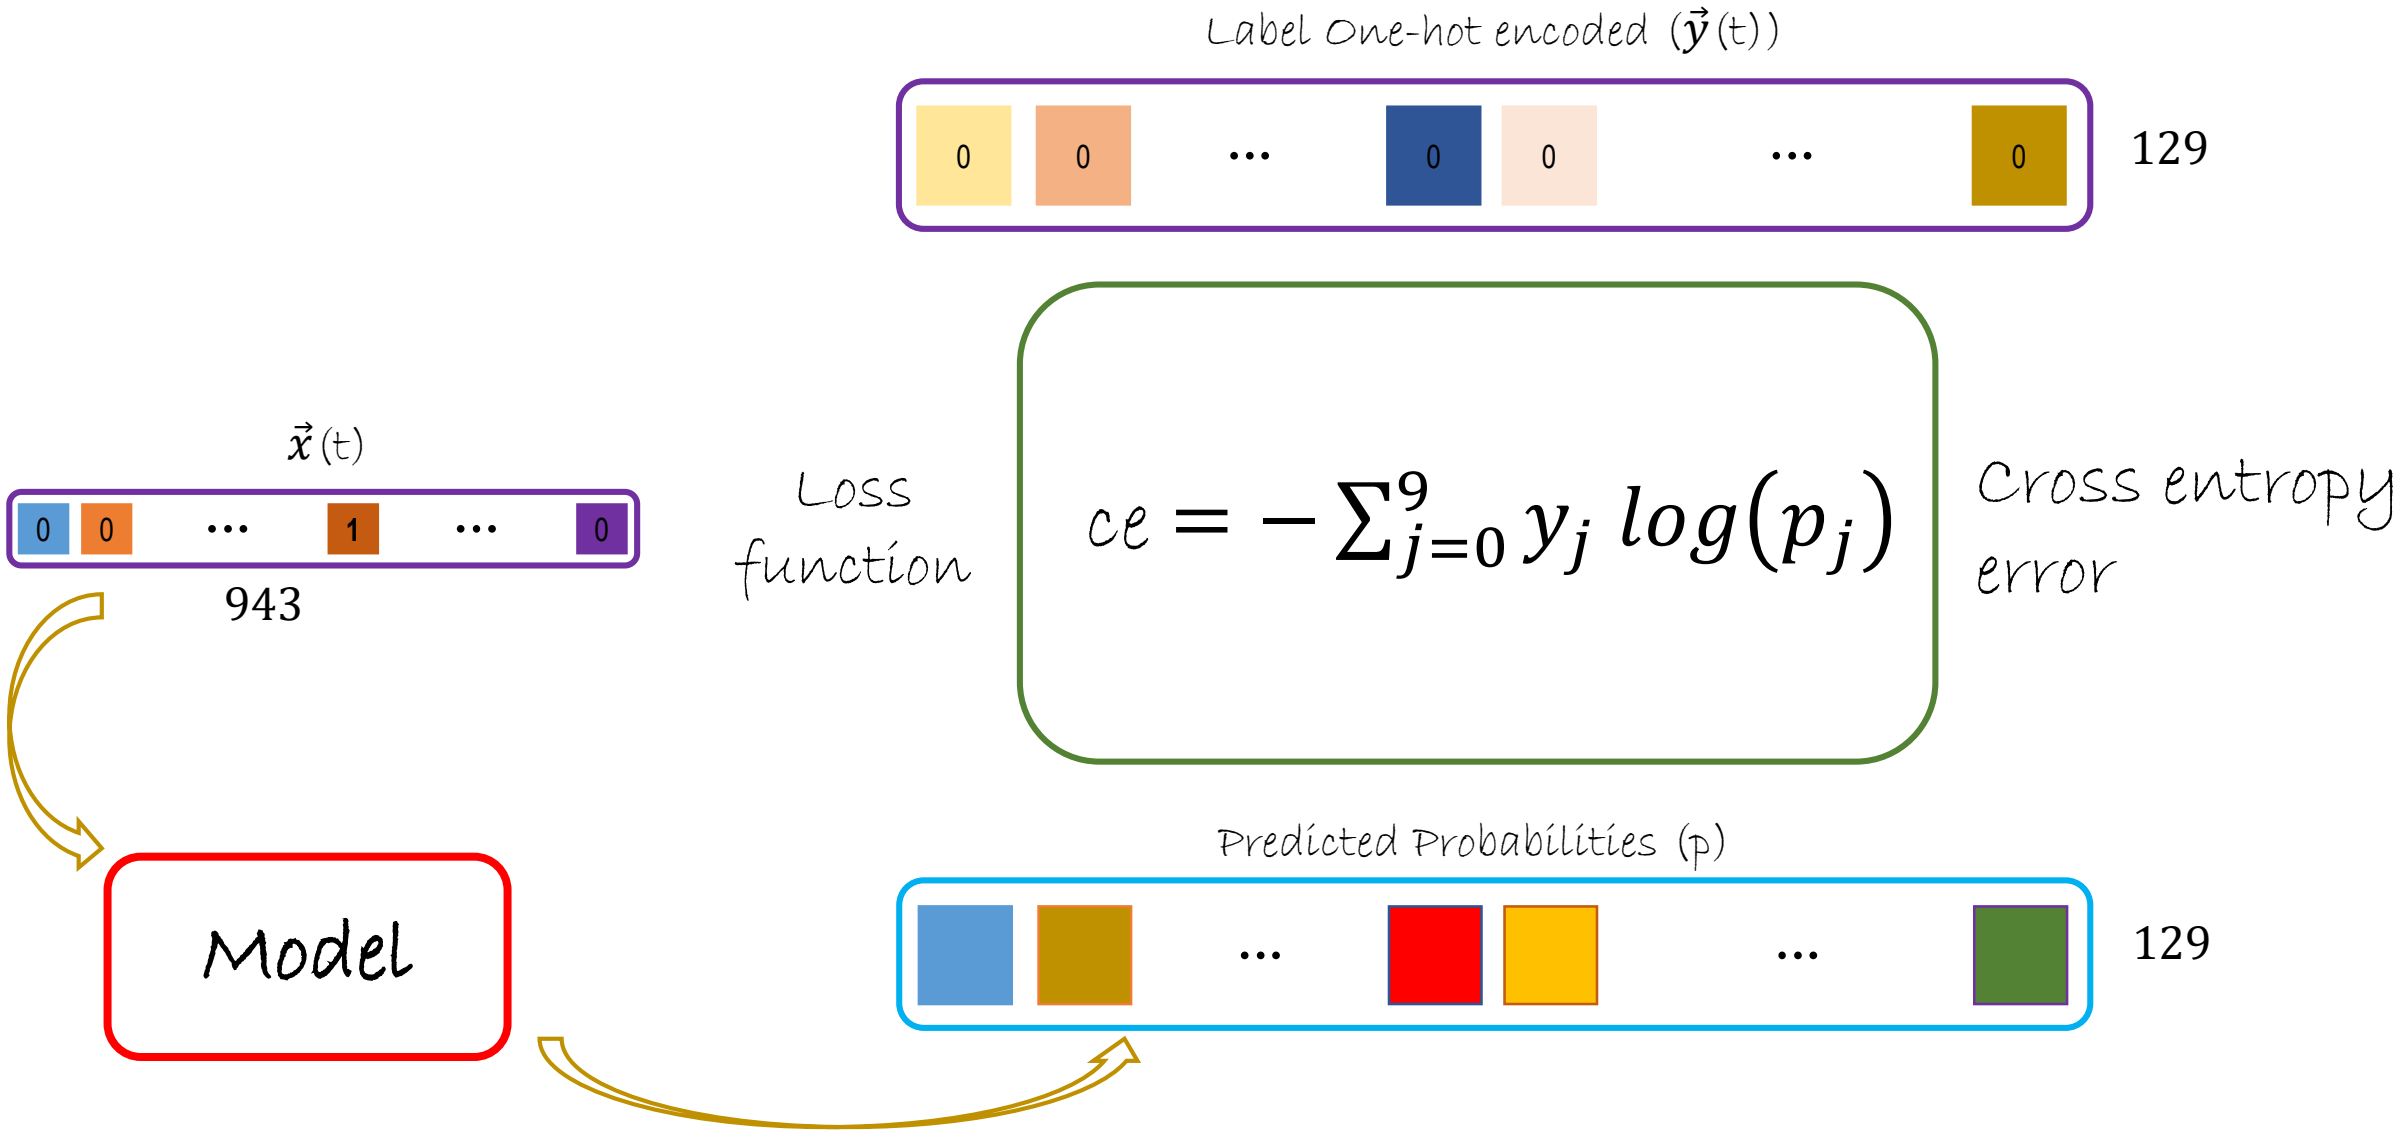

Train / Validation Workflow

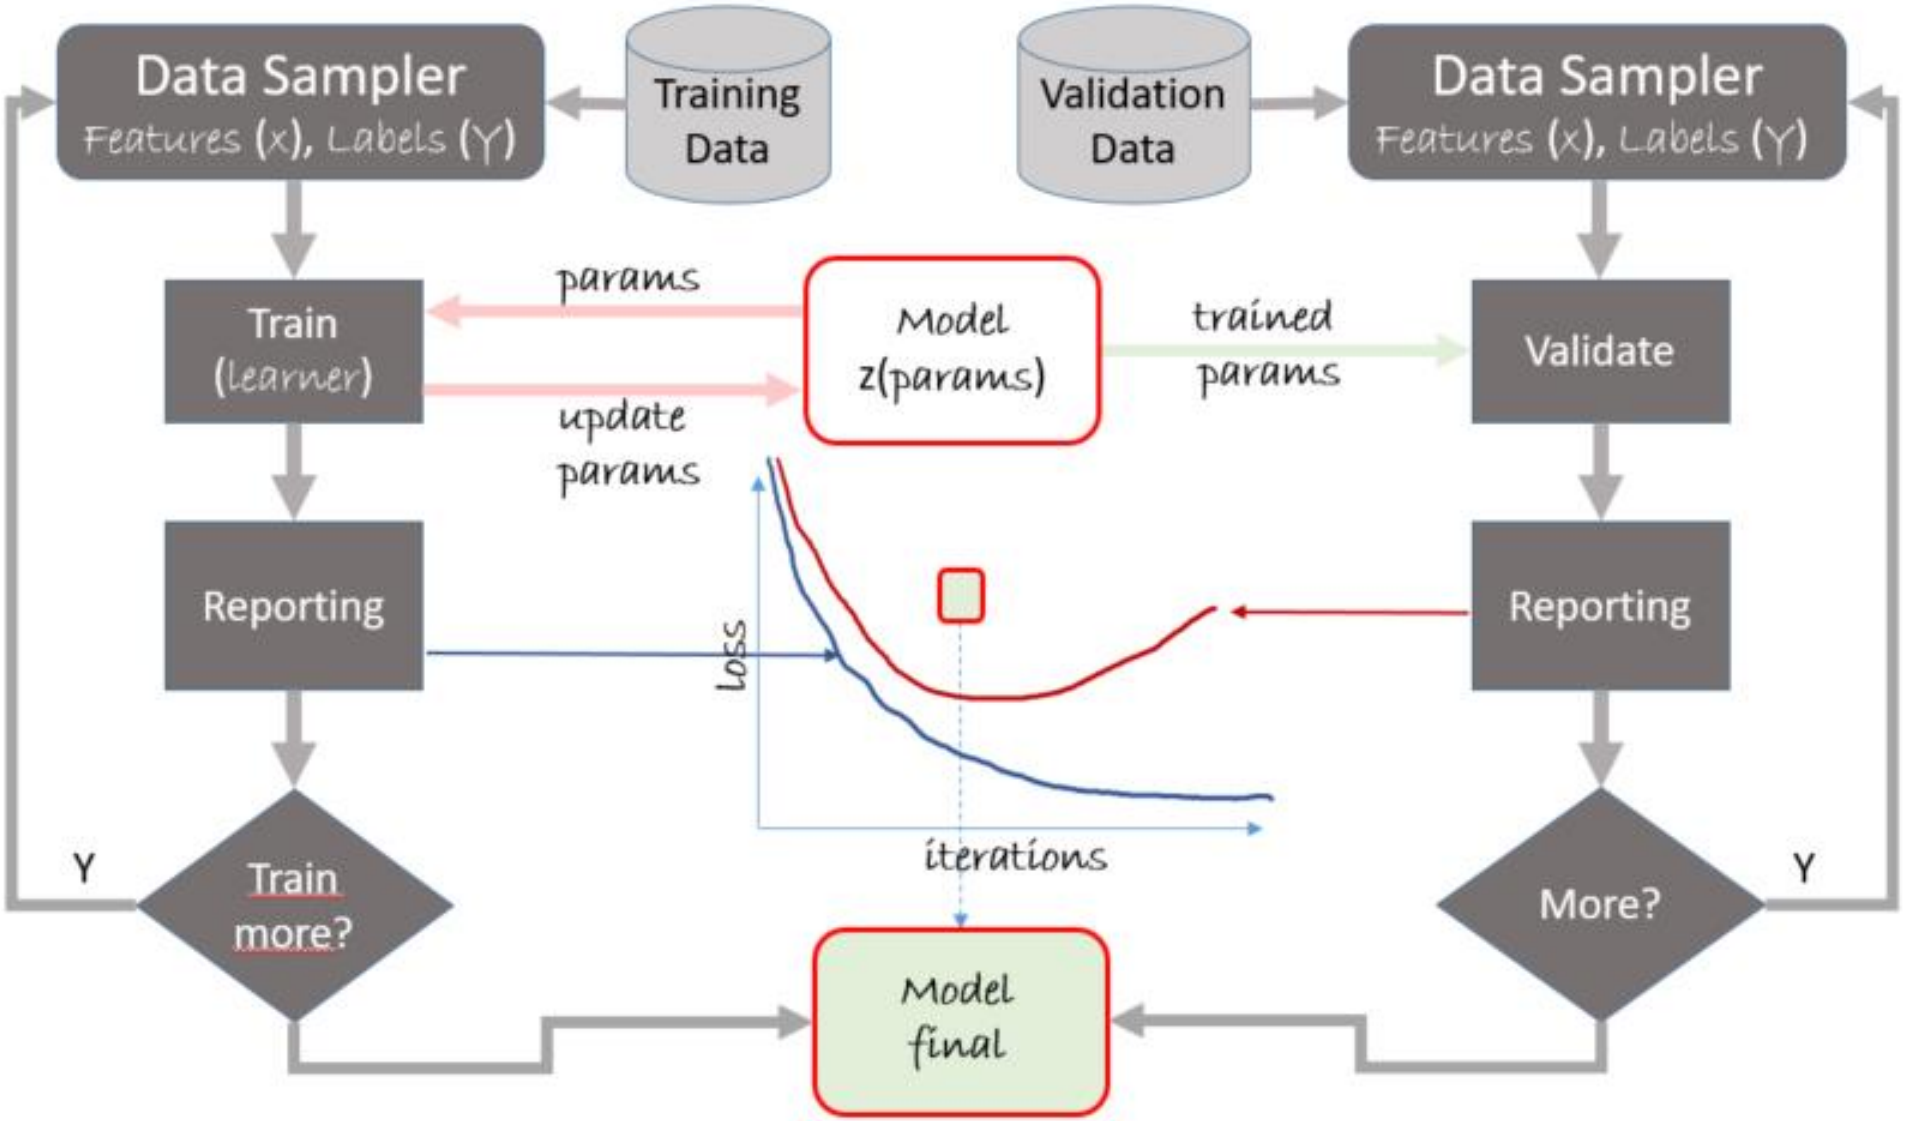

Train Workflow

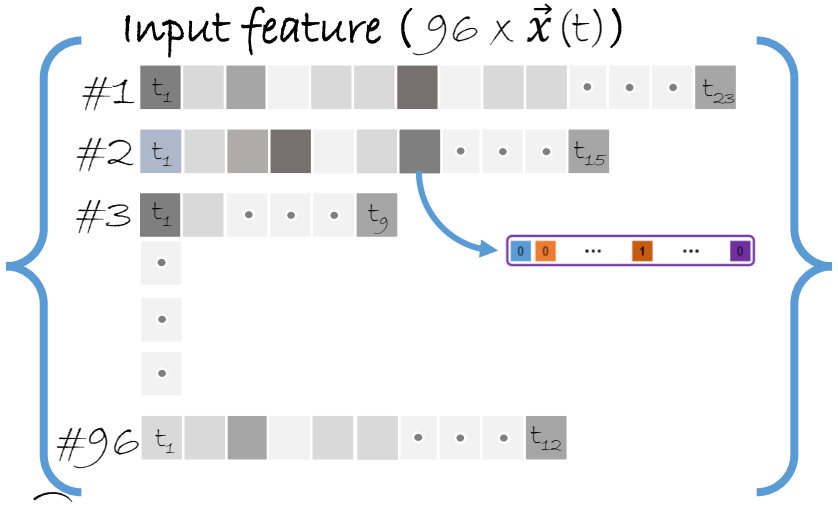


```
z = model():
        return
            Sequential([
                Embedding(emb_dim=150),
                Recurrence(LSTM(hidden_dim=300),
                           go_backwards=False),
                Dense(num_labels = 129)
            ])
```

ATIS Train

96 samples (mini-batch)

One-hot encoded Label (Y: 96 x 129/sample Or word in sequence)

Loss

```
cross_entropy_with_softmax(z,Y)
```

Error

```
classification_error(z,Y)
```

```
Trainer(model, (loss, error), learner)
```

```
Trainer.train_minibatch({X, Y})
```

Learner

Adam, adagrad etc, are solvers to estimate

Test workflow

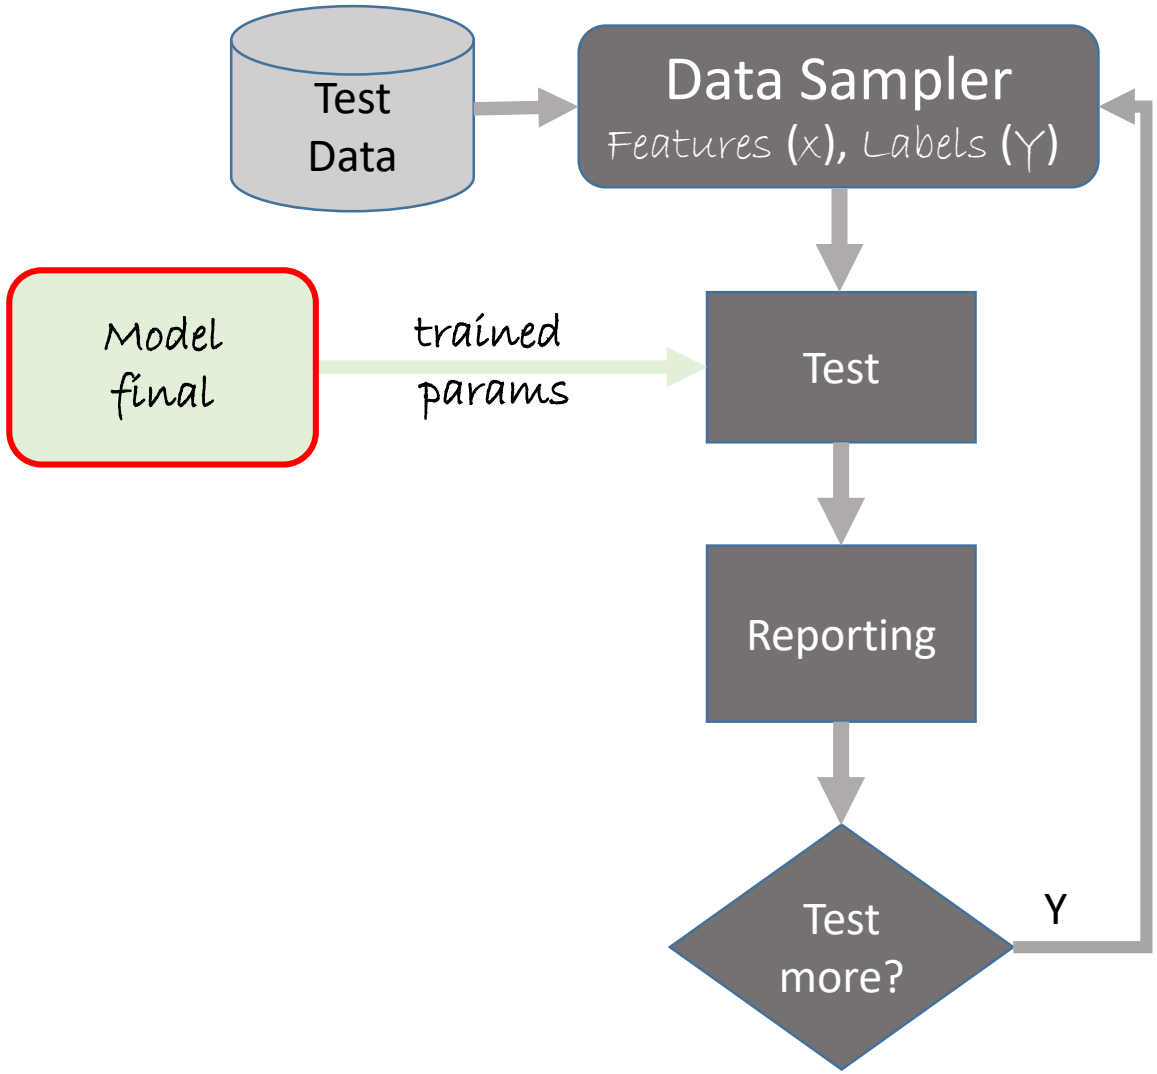

Test workflow


```
z = model():
      return
        Sequential([
            Embedding(emb_dim=150),
            Recurrence(LSTM(hidden_dim=300),
                       go_backwards=False),
            Dense(num_labels = 129)
        ])
```

32 samples (mini-batch)

One-hot encoded Label (Y: 32 x 129/sample Or word in sequence)

```
Trainer.test_minibatch({X, Y})
```

Returns the classification error as % incorrectly labeled tokens.

Prediction workflow

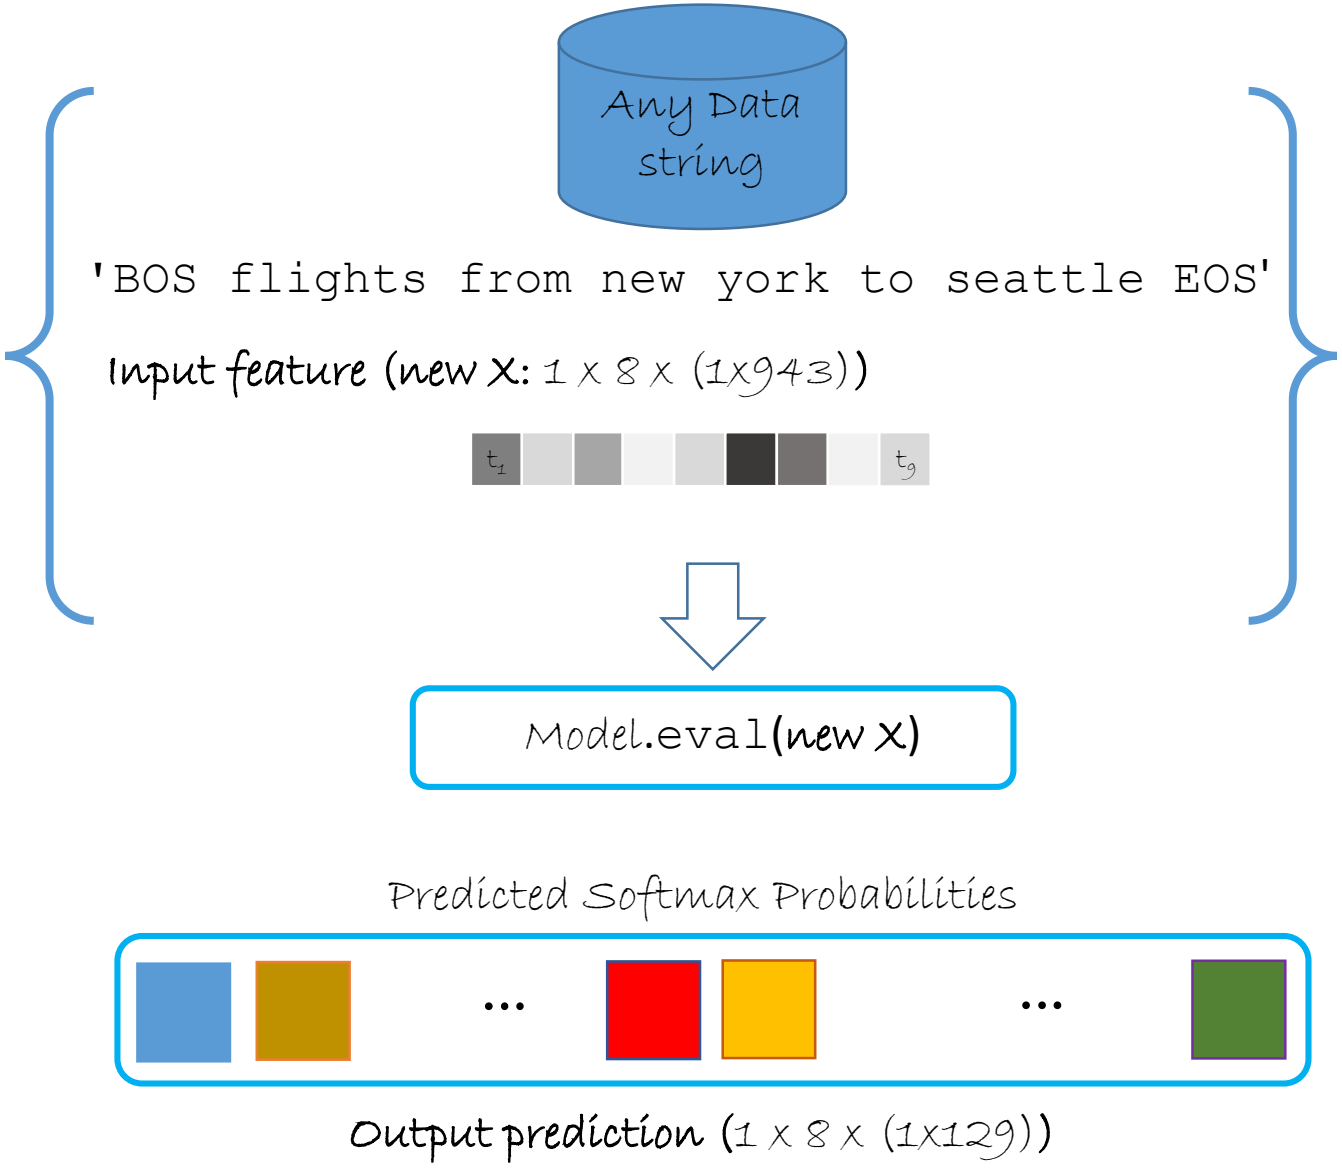

Sequences (many to many)

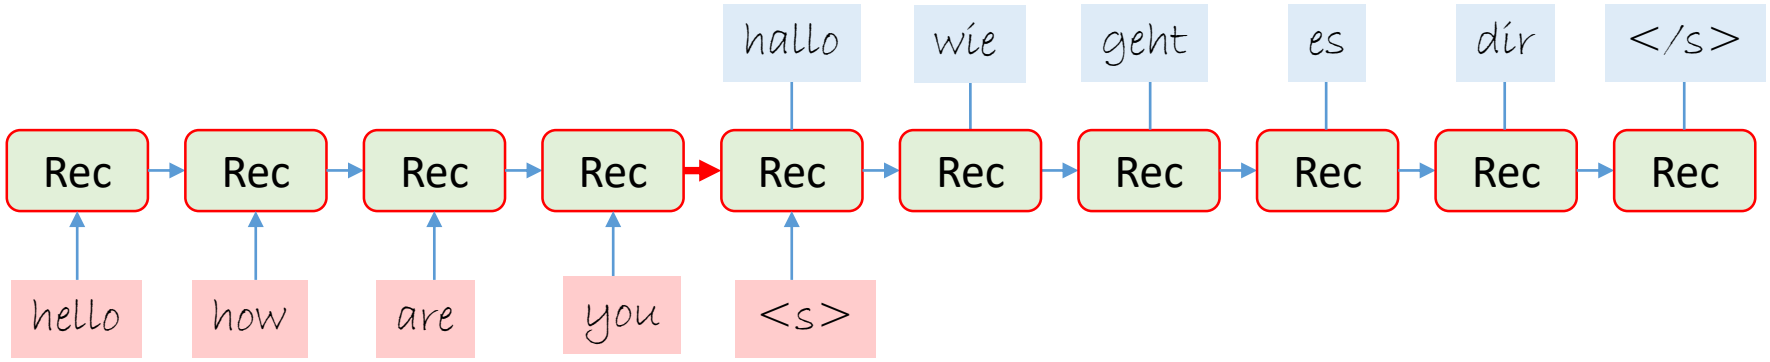

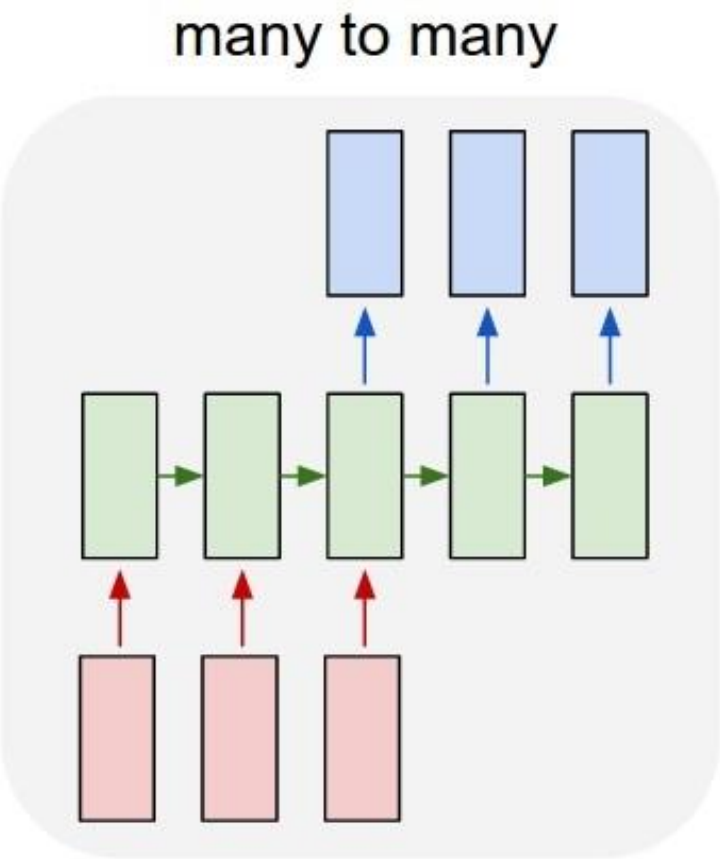

http://karpathy.github.io/2015/05/21/rnn-effectiveness/

Sequences (one to many)

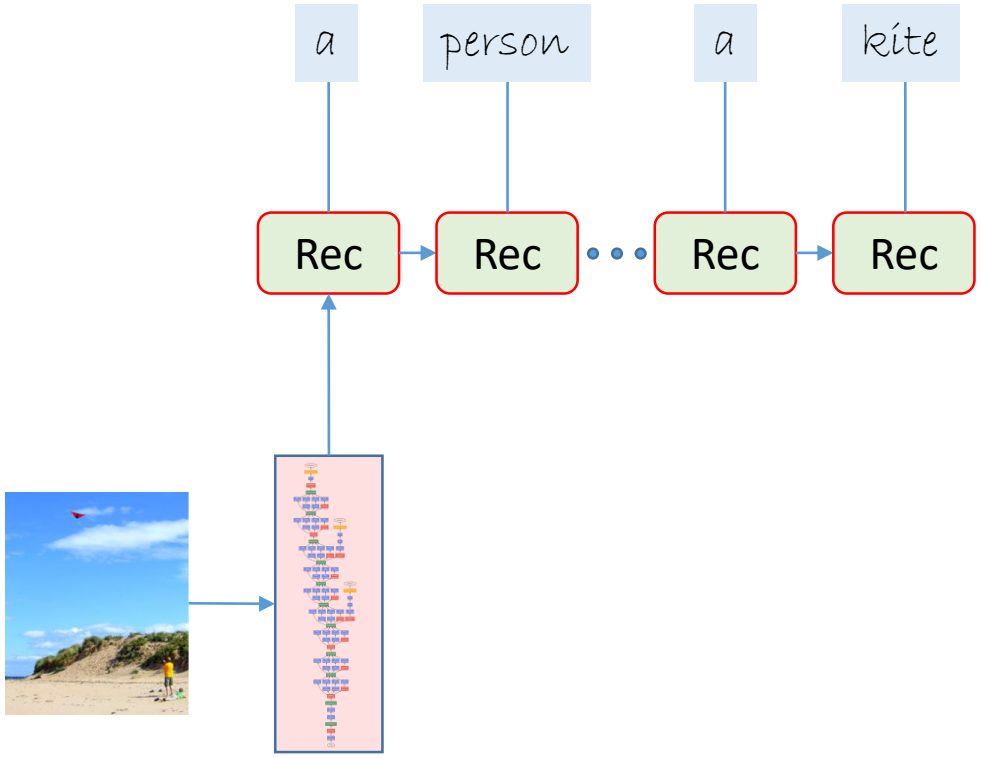

A person on a beach flying a kite.

A black and white photo of a train on a train track.

A person skiing down a snow covered slope.

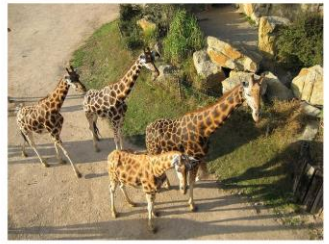

A group of giraffe standing next to each other.

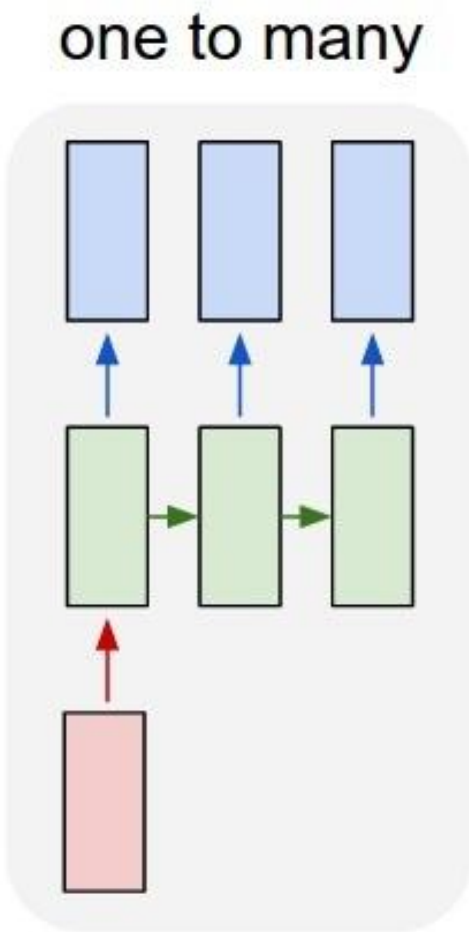

http://karpathy.github.io/2015/05/21/rnn-effectiveness/

Vinyals et al (https://arxiv.org/abs/1411.4555)

Conclusion

Deep learning concepts

- Loss functions, Mini-batch
- Activation functions
- Convolution, Pooling
- Recurrence, LSTM, Dropout, Embeddings

Deep neural networks models

- Multi-class logistic regression
- Multi-layered perceptron
- Convolutional neural networks
- Recurrent networks with LSTM
- Recurrent networks with LSTM and word embeddings

Train-Test-Predict using DNN models